

```
# NAME :- AKULA HEMA VENKATA SRIRAM
# ROLL NO :- 04
# REGISTRATION NO :- 12210461
# SECTION :- K22BW
# COURSE CODE :- INT-354

#
-----
-----
-----
# PROJECT TITLE :- Analyzing Machine Learning Models for Human
Activity Recognition: A Comparative Study
#
-----
-----
-----
```

Reading The Dataset

```
import pandas as pd
data=pd.read_csv('accelerometer_gyro_mobile_phone_dataset.csv')
data.head()
```

	accX	accY	accZ	gyroX	gyroY	gyroZ
timestamp \						
0	-0.496517	3.785628	8.954828	-0.142849	-0.126159	-0.022539
34:22.9						
1	-0.462388	3.869603	9.281898	0.084349	0.096695	0.092130
34:23.0						
2	-0.296084	3.820505	8.930728	0.061763	0.051543	0.071287
34:23.1						
3	-0.469723	3.890110	8.744067	0.007641	0.028679	0.109433
34:23.2						
4	-0.472418	4.109105	8.941207	-0.123640	0.099057	0.051943
34:23.3						

	Activity
0	1
1	1
2	1
3	1
4	1

Data Exploration and Preprocessing

```
# Display the first few rows of the dataset
print("First few rows of the dataset:")
print(data.head())
```

```
# Check the shape of the dataset
print("\nShape of the dataset:")
print(data.shape)
```

First few rows of the dataset:

	accX	accY	accZ	gyroX	gyroY	gyroZ
timestamp \						
0	-0.496517	3.785628	8.954828	-0.142849	-0.126159	-0.022539
34:22.9						
1	-0.462388	3.869603	9.281898	0.084349	0.096695	0.092130
34:23.0						
2	-0.296084	3.820505	8.930728	0.061763	0.051543	0.071287
34:23.1						
3	-0.469723	3.890110	8.744067	0.007641	0.028679	0.109433
34:23.2						
4	-0.472418	4.109105	8.941207	-0.123640	0.099057	0.051943
34:23.3						

	Activity
0	1
1	1
2	1
3	1
4	1

Shape of the dataset:
(31991, 8)

Check For Missing Values

```
print("Missing values:")
print(data.isnull().sum())
```

Missing values:

accX	0
accY	0
accZ	0
gyroX	0
gyroY	0
gyroZ	0
timestamp	0
Activity	0

dtype: int64

```
# Check unique values in the 'activity' column
```

```
print("Unique activities:")
print(data['Activity'].unique())
```

Unique activities:
[1 0]

```
# Summary statistics of numerical columns
```

```
print("Summary statistics:")
```

```
print(data.describe())
```

```
Summary statistics:
```

	accX	accY	accZ	gyroX	
gyroY \					
count	31991.000000	31991.000000	31991.000000	31991.000000	
mean	0.023825	2.153858	9.537909	-0.004493	-
std	0.741396	1.085466	2.056358	0.307643	
min	-3.673361	-4.386029	4.296066	-1.470421	-
25%	-0.472193	1.413062	7.794217	-0.149783	-
50%	-0.024998	2.119143	9.406739	0.022301	-
75%	0.477208	2.928435	11.158845	0.177978	
max	4.678671	6.377039	17.591568	1.332722	

	gyroZ	Activity
count	31991.000000	31991.000000
mean	-0.007021	0.982151
std	0.266120	0.132404
min	-1.894102	0.000000
25%	-0.154126	1.000000
50%	0.005261	1.000000
75%	0.152061	1.000000
max	1.482268	1.000000

Scaling

```
import pandas as pd
```

```
from sklearn.preprocessing import StandardScaler
```

```
# Scaling numerical features
```

```
# We'll scale the accelerometer and gyroscope signals using  
StandardScaler
```

```
scaler = StandardScaler()
```

```
numerical_cols = [  
    'accX', 'accY', 'accZ',  
    'gyroX', 'gyroY', 'gyroZ'  
]
```

```
data[numerical_cols] = scaler.fit_transform(data[numerical_cols])
```

```
# Display the first few rows of the updated dataset with new features
and preprocessed data
print("Updated dataset after preprocessing:")
print(data.head())
```

Updated dataset after preprocessing:

	accX	accY	accZ	gyroX	gyroY	gyroZ
timestamp \						
0	-0.701852	1.503313	-0.283555	-0.449734	-0.447050	-0.058313
34:22.9						
1	-0.655818	1.580678	-0.124499	0.288789	0.447239	0.372586
34:23.0						
2	-0.431503	1.535445	-0.295275	0.215371	0.266049	0.294263
34:23.1						
3	-0.665712	1.599570	-0.386049	0.039444	0.174298	0.437607
34:23.2						
4	-0.669347	1.801326	-0.290179	-0.387294	0.456717	0.221573
34:23.3						

	Activity
0	1
1	1
2	1
3	1
4	1

Check if the data is imbalanced

```
print(100*data['Activity'].value_counts()/len(data['Activity']))
print(data['Activity'].value_counts())
```

```
Activity
1    98.215123
0     1.784877
Name: count, dtype: float64
Activity
1    31420
0     571
Name: count, dtype: int64
```

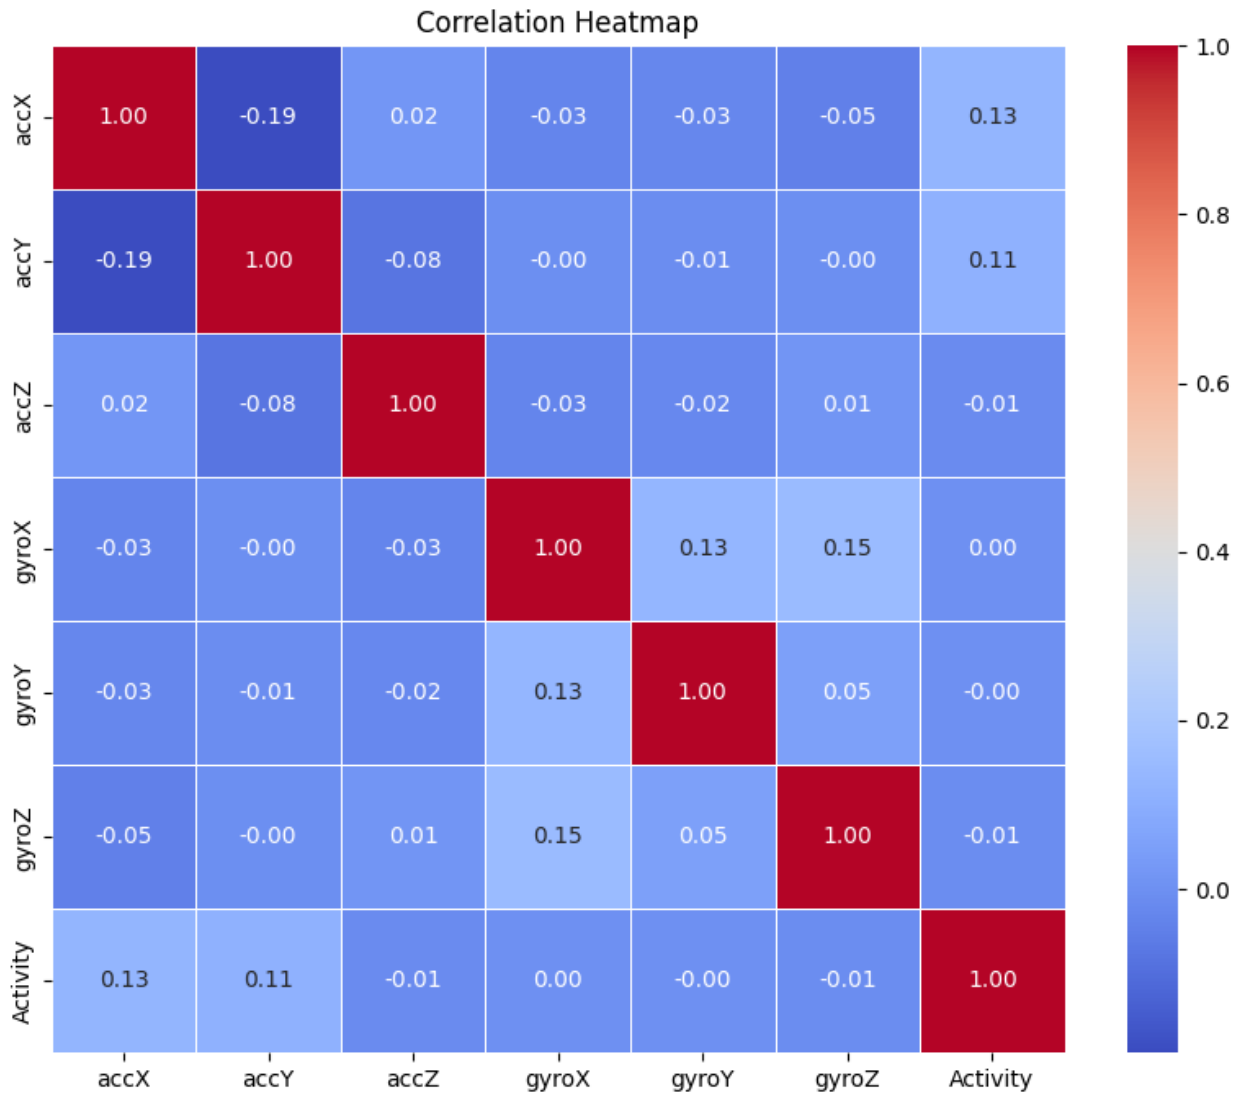
Plotting Heatmap

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Drop the 'timestamp' column
data = data.drop(columns=['timestamp'])
```

```
# Compute the correlation matrix
correlation_matrix = data.corr()

# Plot the heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
            fmt=".2f",linewidths=0.5)
plt.title('Correlation Heatmap')
plt.show()
```



```
df=data.drop(columns=['gyroX','gyroY'])
df.head()
```

	accX	accY	accZ	gyroZ	Activity
0	-0.701852	1.503313	-0.283555	-0.058313	1
1	-0.655818	1.580678	-0.124499	0.372586	1

2	-0.431503	1.535445	-0.295275	0.294263	1
3	-0.665712	1.599570	-0.386049	0.437607	1
4	-0.669347	1.801326	-0.290179	0.221573	1

Model Training & Evaluation

Splitting data and target

```
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix

# Split the data into features (X) and target (y)
X = df.drop(columns=['Activity']) # Features
y = df['Activity'] # Target

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix

# Define the Random Forest Classifier with default hyperparameters
rf_classifier = RandomForestClassifier(random_state=42) # Set
random_state for reproducibility

# Train the model on the training data
rf_classifier.fit(X_train, y_train)

# Make predictions on the testing data
y_pred_rf = rf_classifier.predict(X_test)

# Evaluate the model's performance
print("Classification Report:")
print(classification_report(y_test, y_pred_rf))

print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred_rf))
```

Classification Report:

	precision	recall	f1-score	support
0	0.69	0.19	0.30	106
1	0.99	1.00	0.99	6293
accuracy			0.99	6399
macro avg	0.84	0.59	0.64	6399
weighted avg	0.98	0.99	0.98	6399

Confusion Matrix:

```
[[ 20 86]
 [ 9 6284]]
```

OverSampling

```
from imblearn.combine import SMOTEENN
sm=SMOTEENN()
X_resampled,Y_resampled=sm.fit_resample(X,y)
Xr_train,Xr_test,Yr_train,Yr_test=train_test_split(X_resampled,Y_resampled,test_size=0.2,random_state=42)
print(y.value_counts())
print("\nAfter Resampling")
print(Y_resampled.value_counts())
```

```
Activity
1    31420
0      571
Name: count, dtype: int64
```

```
After Resampling
Activity
0    30908
1    28972
Name: count, dtype: int64
```

WITHOUT HYPERPARAMETER TUNING

Random Forest Classifier

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix

# Define the Random Forest Classifier with default hyperparameters
rf_classifier = RandomForestClassifier(random_state=42) # Set random_state for reproducibility

# Train the model on the training data
rf_classifier.fit(Xr_train, Yr_train)

# Make predictions on the testing data
Y_pred_rf = rf_classifier.predict(Xr_test)

# Evaluate the model's performance
print("Classification Report:")
print(classification_report(Yr_test, Y_pred_rf))

print("\nConfusion Matrix:")
print(confusion_matrix(Yr_test, Y_pred_rf))
```

```

Classification Report:
              precision    recall  f1-score   support

     0       0.98        1.00        0.99        6148
     1       1.00        0.98        0.99        5828

 accuracy          0.99          0.99          0.99        11976
 macro avg          0.99          0.99          0.99        11976
 weighted avg       0.99          0.99          0.99        11976

Confusion Matrix:
[[6125   23]
 [  98 5730]]

```

Decision Tree Classifier

```

from sklearn.pipeline import Pipeline
from sklearn.tree import DecisionTreeClassifier

# Define pipeline with Decision Tree Classifier (default
hyperparameters)
pipeline = Pipeline([
    ('classifier', DecisionTreeClassifier(random_state=42))
])

# Train the pipeline on the training data
pipeline.fit(Xr_train, Yr_train)

# Make predictions on the testing data
Y_pred_dt = pipeline.predict(Xr_test)

# Evaluate the model's performance
print("\nClassification Report:")
print(classification_report(Yr_test, Y_pred_dt))

print("\nConfusion Matrix:")
print(confusion_matrix(Yr_test, Y_pred_dt))

```

```

Classification Report:
              precision    recall  f1-score   support

     0       0.98        0.99        0.98        6148
     1       0.99        0.98        0.98        5828

 accuracy          0.98          0.98          0.98        11976
 macro avg          0.98          0.98          0.98        11976
 weighted avg       0.98          0.98          0.98        11976

```


Confusion Matrix:

```
[[6067   81]
 [ 139 5689]]
```

Logistic Regression

```
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression

# Define pipeline with Logistic Regression (default hyperparameters)
pipeline = Pipeline([
    ('classifier', LogisticRegression(random_state=42))
])

# Train the pipeline on the training data
pipeline.fit(Xr_train, Yr_train)

# Make predictions on the testing data
Y_pred_lr = pipeline.predict(Xr_test)

# Evaluate the performance of the Logistic Regression model
print("\nClassification Report for Logistic Regression:")
print(classification_report(Yr_test, Y_pred_lr))

print("\nConfusion Matrix for Logistic Regression:")
print(confusion_matrix(Yr_test, Y_pred_lr))
```

Classification Report for Logistic Regression:

	precision	recall	f1-score	support
0	0.85	0.93	0.89	6148
1	0.92	0.83	0.87	5828
accuracy			0.88	11976
macro avg	0.89	0.88	0.88	11976
weighted avg	0.89	0.88	0.88	11976

Confusion Matrix for Logistic Regression:

```
[[5745  403]
 [ 993 4835]]
```

K Neighbors Classifier

```
from sklearn.neighbors import KNeighborsClassifier

# Define the pipeline with KNeighborsClassifier (default
hyperparameters)
```

```

pipeline = Pipeline([
    ('classifier', KNeighborsClassifier())
])

# Train the pipeline on the training data
pipeline.fit(Xr_train, Yr_train)

# Make predictions on the testing data
Y_pred_knn = pipeline.predict(Xr_test)

# Evaluate the model's performance
print("\nClassification Report:")
print(classification_report(Yr_test, Y_pred_knn))

print("\nConfusion Matrix:")
print(confusion_matrix(Yr_test, Y_pred_knn))

```

Classification Report:

	precision	recall	f1-score	support
0	0.98	1.00	0.99	6148
1	1.00	0.98	0.99	5828
accuracy			0.99	11976
macro avg	0.99	0.99	0.99	11976
weighted avg	0.99	0.99	0.99	11976

Confusion Matrix:

```

[[6139    9]
 [ 101 5727]]

```

WITH HYPERPARAMETER TUNING - GRIDSEARCH

Decision tree classifier

```

from sklearn.model_selection import GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.tree import DecisionTreeClassifier

# Define pipeline with Decision Tree Classifier
pipeline = Pipeline([
    ('classifier', DecisionTreeClassifier(random_state=42))
])

# Define hyperparameters grid for grid search
param_grid = {
    'classifier__max_depth': [None, 10, 20],
    'classifier__min_samples_split': [2, 5, 10],

```

```

    'classifier__min_samples_leaf': [1, 2, 4]
}

# Perform hyperparameter tuning using GridSearchCV
grid_search = GridSearchCV(estimator=pipeline, param_grid=param_grid,
cv=5, scoring='accuracy')
grid_search.fit(Xr_train, Yr_train)

# Get the best parameters
best_params = grid_search.best_params_
print("Best Parameters:", best_params)

# Predict on the testing set using the best model
best_classifier = grid_search.best_estimator_
Y_pred_dt_grid = best_classifier.predict(Xr_test)

# Evaluate the model
print("\nClassification Report:")
print(classification_report(Yr_test, Y_pred_dt))

print("\nConfusion Matrix:")
print(confusion_matrix(Yr_test, Y_pred_dt))

Best Parameters: {'classifier__max_depth': None,
'classifier__min_samples_leaf': 1, 'classifier__min_samples_split': 2}

Classification Report:

```

	precision	recall	f1-score	support
0	0.98	0.99	0.98	6148
1	0.99	0.98	0.98	5828
accuracy			0.98	11976
macro avg	0.98	0.98	0.98	11976
weighted avg	0.98	0.98	0.98	11976

```

Confusion Matrix:
[[6067  81]
 [ 139 5689]]

```

Logistic Regression

```

pipeline = Pipeline([
    ('classifier', LogisticRegression(random_state=42)) # Logistic
Regression classifier
])

# Define hyperparameter grid for GridSearchCV
param_grid = {

```

```

    'classifier__C': [0.001, 0.01, 0.1, 1, 10], # Specify a list of
values for C
    'classifier__solver': ['lbfgs', 'liblinear'] # Solvers to try
}

# Perform hyperparameter tuning using GridSearchCV
grid_search = GridSearchCV(estimator=pipe, param_grid=param_grid,
cv=5, scoring='accuracy')
grid_search.fit(Xr_train, Yr_train)

# Get the best parameters
best_params = grid_search.best_params_
print("Best Parameters:", best_params)

# Predict on the testing set using the best model
best_classifier = grid_search.best_estimator_
Y_pred_lr_grid = best_classifier.predict(Xr_test)

# Evaluate the performance of the Logistic Regression model
print("\nClassification Report for Logistic Regression:")
print(classification_report(Yr_test, Y_pred_lr))

print("\nConfusion Matrix for Logistic Regression:")
print(confusion_matrix(Yr_test, Y_pred_lr))

```

Best Parameters: {'classifier__C': 10, 'classifier__solver': 'lbfgs'}

Classification Report for Logistic Regression:

	precision	recall	f1-score	support
0	0.85	0.93	0.89	6148
1	0.92	0.83	0.87	5828
accuracy			0.88	11976
macro avg	0.89	0.88	0.88	11976
weighted avg	0.89	0.88	0.88	11976

Confusion Matrix for Logistic Regression:

```
[[5745  403]
 [ 993 4835]]
```

WITH HYPERPARAMETER TUNING RANDOMIZEDSEARCHCV

Decision Tree Classifier

```

from sklearn.model_selection import RandomizedSearchCV
from sklearn.pipeline import Pipeline
from sklearn.tree import DecisionTreeClassifier
from scipy.stats import uniform, randint

```

```

# Define pipeline with Decision Tree Classifier
pipeline = Pipeline([
    ('classifier', DecisionTreeClassifier(random_state=42))
])

# Define hyperparameter distributions for RandomizedSearchCV
# Define hyperparameter distributions for RandomizedSearchCV
param_dist = {
    'classifier__criterion': ['gini', 'entropy'], # Options for
criterion
    'classifier__max_depth': randint(2, 20),      # Integer values
between 2 and 20
    'classifier__min_samples_split': randint(2, 10), # Integer values
between 2 and 10
    'classifier__min_samples_leaf': randint(1, 4)  # Integer values
between 1 and 4
}

# Perform hyperparameter tuning using RandomizedSearchCV
random_dt = RandomizedSearchCV(estimator=pipeline,
param_distributions=param_dist, n_iter=10, cv=5, scoring='accuracy',
random_state=42)
random_dt.fit(Xr_train, Yr_train)

# Get the best parameters
best_params = random_dt.best_params_
print("Best Parameters:", best_params)

# Predict on the testing set using the best model
best_classifier = random_dt.best_estimator_
Y_pred_dt_rand = best_classifier.predict(Xr_test)

# Evaluate the model
print("\nClassification Report:")
print(classification_report(Yr_test, Y_pred_dt))

print("\nConfusion Matrix:")
print(confusion_matrix(Yr_test, Y_pred_dt))

Best Parameters: {'classifier__criterion': 'entropy',
'classifier__max_depth': 16, 'classifier__min_samples_leaf': 2,
'classifier__min_samples_split': 4}

```

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.99	0.98	6148
1	0.99	0.98	0.98	5828

accuracy			0.98	11976
macro avg	0.98	0.98	0.98	11976
weighted avg	0.98	0.98	0.98	11976

Confusion Matrix:

```
[[6067  81]
 [ 139 5689]]
```

Logistic Regression

```
from sklearn.model_selection import RandomizedSearchCV
from sklearn.linear_model import LogisticRegression
from scipy.stats import uniform

# Define the pipeline
pipeline = Pipeline([
    ('classifier', LogisticRegression(random_state=42))
])

# Define hyperparameters to search
param_dist = {
    'classifier__C': uniform(0.001, 10), # Inverse regularization
    strength
    'classifier__solver': ['lbfgs', 'liblinear'] # Solvers to try
}

# Perform hyperparameter tuning using RandomizedSearchCV
random_lr = RandomizedSearchCV(estimator=pipeline,
    param_distributions=param_dist, n_iter=10, cv=5, scoring='accuracy',
    random_state=42)
random_lr.fit(Xr_train, Yr_train)

# Get the best parameters
best_params = random_lr.best_params_
print("Best Parameters:", best_params)

# Predict on the testing set using the best model
best_classifier = random_lr.best_estimator_
Y_pred_lr_rand = best_classifier.predict(Xr_test)

# Evaluate the performance of the Logistic Regression model
print("\nClassification Report for Logistic Regression:")
print(classification_report(Yr_test, Y_pred_lr))

print("\nConfusion Matrix for Logistic Regression:")
print(confusion_matrix(Yr_test, Y_pred_lr))

Best Parameters: {'classifier__C': 3.746401188473625,
'classifier__solver': 'lbfgs'}
```

Classification Report for Logistic Regression:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.85	0.93	0.89	6148
1	0.92	0.83	0.87	5828
accuracy			0.88	11976
macro avg	0.89	0.88	0.88	11976
weighted avg	0.89	0.88	0.88	11976

Confusion Matrix for Logistic Regression:

```
[[5745  403]
 [ 993 4835]]
```

KNN

```
from sklearn.neighbors import KNeighborsClassifier
from scipy.stats import randint

# Define the pipeline
pipeline = Pipeline([ ('classifier', KNeighborsClassifier())])

param_dist = {
    'classifier__n_neighbors': randint(1, 20), # Number of neighbors
    'classifier__p': [1, 2] # Distance metric: 1 for Manhattan, 2 for
    Euclidean
}

# Perform hyperparameter tuning using RandomizedSearchCV
random_knn = RandomizedSearchCV(estimator=pipeline,
    param_distributions=param_dist, n_iter=10, cv=5, scoring='accuracy',
    random_state=42)
random_knn.fit(Xr_train, Yr_train)

# Get the best parameters
best_params = random_knn.best_params_
print("Best Parameters:", best_params)

# Predict on the testing set using the best model
best_classifier = random_knn.best_estimator_
Y_pred_knn_rand = best_classifier.predict(Xr_test)

# Evaluate the model
print("\nClassification Report:")
print(classification_report(Yr_test, Y_pred_knn))

print("\nConfusion Matrix:")
print(confusion_matrix(Yr_test, Y_pred_knn))
```

Best Parameters: {'classifier__n_neighbors': 3, 'classifier__p': 2}

Classification Report:

	precision	recall	f1-score	support
0	0.98	1.00	0.99	6148
1	1.00	0.98	0.99	5828
accuracy			0.99	11976
macro avg	0.99	0.99	0.99	11976
weighted avg	0.99	0.99	0.99	11976

Confusion Matrix:

```
[[6139    9]
 [ 101 5727]]
```

Comparing Scores

```
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score

# Calculate metrics for each model
dt_metrics = [accuracy_score(Yr_test, Y_pred_dt),
               precision_score(Yr_test, Y_pred_dt, average='weighted'),
               recall_score(Yr_test, Y_pred_dt, average='weighted'),
               f1_score(Yr_test, Y_pred_dt, average='weighted')]

rf_metrics = [accuracy_score(Yr_test, Y_pred_rf),
               precision_score(Yr_test, Y_pred_rf, average='weighted'),
               recall_score(Yr_test, Y_pred_rf, average='weighted'),
               f1_score(Yr_test, Y_pred_rf, average='weighted')]

knn_metrics = [accuracy_score(Yr_test, Y_pred_knn),
               precision_score(Yr_test, Y_pred_knn,
                               average='weighted'),
               recall_score(Yr_test, Y_pred_knn, average='weighted'),
               f1_score(Yr_test, Y_pred_knn, average='weighted')]

lr_metrics = [accuracy_score(Yr_test, Y_pred_lr),
               precision_score(Yr_test, Y_pred_lr, average='weighted'),
               recall_score(Yr_test, Y_pred_lr, average='weighted'),
               f1_score(Yr_test, Y_pred_lr, average='weighted')]

# Metrics dataframe
metrics_df = pd.DataFrame([dt_metrics, rf_metrics, knn_metrics,
                             lr_metrics],
                           columns=['Accuracy', 'Precision', 'Recall',
                                   'F1-Score'],
                           index=['DecisionTreeClassifier',
                                   'RandomForestClassifier', 'KNN', 'LogisticRegression'])
```



```

# Print the metrics
print("Metrics for DecisionTreeClassifier:")
print(metrics_df.loc['DecisionTreeClassifier'])
print("\nMetrics for RandomForestClassifier:")
print(metrics_df.loc['RandomForestClassifier'])
print("\nMetrics for KNN:")
print(metrics_df.loc['KNN'])
print("\nMetrics for Logistic Regression:")
print(metrics_df.loc['LogisticRegression'])

# Print the model with the best F1-score
best_model = metrics_df.idxmax()['Accuracy']
print("\nBest Model for the Project based on Accuracy:", best_model)

```

```

Metrics for DecisionTreeClassifier:
Accuracy      0.981630
Precision      0.981670
Recall         0.981630
F1-Score       0.981627
Name: DecisionTreeClassifier, dtype: float64

```

```

Metrics for RandomForestClassifier:
Accuracy      0.989896
Precision      0.989970
Recall         0.989896
F1-Score       0.989894
Name: RandomForestClassifier, dtype: float64

```

```

Metrics for KNN:
Accuracy      0.990815
Precision      0.990927
Recall         0.990815
F1-Score       0.990813
Name: KNN, dtype: float64

```

```

Metrics for Logistic Regression:
Accuracy      0.883434
Precision      0.886904
Recall         0.883434
F1-Score       0.882995
Name: LogisticRegression, dtype: float64

```

```

Best Model for the Project based on Accuracy: KNN

```

```

from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score

```

```

# Calculate metrics for each model
dt_metrics_grid = [accuracy_score(Yr_test, Y_pred_dt_grid),

```

```

        precision_score(Yr_test, Y_pred_dt_grid,
average='weighted'),
        recall_score(Yr_test, Y_pred_dt_grid,
average='weighted'),
        f1_score(Yr_test, Y_pred_dt_grid, average='weighted')])

lr_metrics_grid = [accuracy_score(Yr_test, Y_pred_lr_grid),
        precision_score(Yr_test, Y_pred_lr_grid,
average='weighted'),
        recall_score(Yr_test, Y_pred_lr_grid,
average='weighted'),
        f1_score(Yr_test, Y_pred_lr_grid, average='weighted')]

# Metrics dataframe
metrics_df_grid = pd.DataFrame([dt_metrics_grid,lr_metrics_grid],
                                columns=['Accuracy', 'Precision', 'Recall',
'F1-Score'],

index=['DecisionTreeClassifier_grid','LogisticRegression_grid'])

# Print the metrics
print("Metrics for DecisionTreeClassifier:")
print(metrics_df_grid.loc['DecisionTreeClassifier_grid'])
print("\nMetrics for Logistic Regression:")
print(metrics_df_grid.loc['LogisticRegression_grid'])

# Print the model with the best F1-score
best_model = metrics_df_grid.idxmax()['Accuracy']
print("\nBest Model for the Project based on Accuracy:", best_model)

Metrics for DecisionTreeClassifier:
Accuracy      0.981630
Precision      0.981670
Recall         0.981630
F1-Score       0.981627
Name: DecisionTreeClassifier_grid, dtype: float64

Metrics for Logistic Regression:
Accuracy      0.883434
Precision      0.886904
Recall         0.883434
F1-Score       0.882995
Name: LogisticRegression_grid, dtype: float64

Best Model for the Project based on Accuracy:
DecisionTreeClassifier_grid

from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score

```

```

# Calculate metrics for each model
dt_metrics_rand = [accuracy_score(Yr_test, Y_pred_dt_rand),
                    precision_score(Yr_test, Y_pred_dt_rand,
                                    average='weighted'),
                    recall_score(Yr_test, Y_pred_dt_rand,
                                 average='weighted'),
                    f1_score(Yr_test, Y_pred_dt_rand, average='weighted')]
knn_metrics_rand = [accuracy_score(Yr_test, Y_pred_knn_rand),
                    precision_score(Yr_test, Y_pred_knn_rand,
                                    average='weighted'),
                    recall_score(Yr_test, Y_pred_knn_rand,
                                 average='weighted'),
                    f1_score(Yr_test, Y_pred_knn_rand, average='weighted')]

lr_metrics_rand = [accuracy_score(Yr_test, Y_pred_lr_rand),
                   precision_score(Yr_test, Y_pred_lr_rand,
                                   average='weighted'),
                   recall_score(Yr_test, Y_pred_lr_rand,
                                average='weighted'),
                   f1_score(Yr_test, Y_pred_lr_rand, average='weighted')]

# Metrics dataframe
metrics_df_rand = pd.DataFrame([dt_metrics_rand, knn_metrics_rand,
                                lr_metrics_rand],
                                columns=['Accuracy', 'Precision', 'Recall',
                                         'F1-Score'],
                                index=['DecisionTreeClassifier_rand',
                                       'KNN_rand', 'LogisticRegression_rand'])

# Print the metrics
print("Metrics for DecisionTreeClassifier:")
print(metrics_df_rand.loc['DecisionTreeClassifier_rand'])
print("\nMetrics for KNN:")
print(metrics_df_rand.loc['KNN_rand'])
print("\nMetrics for Logistic Regression:")
print(metrics_df_rand.loc['LogisticRegression_rand'])

# Print the model with the best F1-score
best_model = metrics_df_rand.idxmax()['Accuracy']
print("\nBest Model for the Project based on Accuracy:", best_model)

Metrics for DecisionTreeClassifier:
Accuracy      0.981212
Precision      0.981399
Recall         0.981212
F1-Score       0.981205
Name: DecisionTreeClassifier_rand, dtype: float64

Metrics for KNN:
Accuracy      0.995908
Precision      0.995929

```

```
Recall      0.995908
F1-Score    0.995908
Name: KNN_rand, dtype: float64
```

Metrics for Logistic Regression:

```
Accuracy    0.883434
Precision   0.886904
Recall      0.883434
F1-Score    0.882995
Name: LogisticRegression_rand, dtype: float64
```

Best Model for the Project based on Accuracy: KNN_rand

Accuracy Visualization

RandomizedSearchCv

```
from sklearn.metrics import accuracy_score

# Calculate accuracy for each model
dt_accuracy_rand = accuracy_score(Yr_test, Y_pred_dt_rand)
knn_accuracy_rand = accuracy_score(Yr_test, Y_pred_knn_rand)
lr_accuracy_rand = accuracy_score(Yr_test, Y_pred_lr_rand)


# Create the accuracy_rand list
accuracy_rand = [dt_accuracy_rand, knn_accuracy_rand,
lr_accuracy_rand]

# Model names list
model_names_rand = ['DecisionTreeClassifier', 'KNN',
'LogisticRegression']

import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))
plt.barh(model_names_rand, accuracy_rand, color='coral', alpha=0.7)
plt.xlabel('Accuracy')
plt.ylabel('Model')
plt.title('Model Accuracy (Random Hyperparameters)')
plt.gca().invert_yaxis() # Invert y-axis for readability (highest
accuracy on top)
plt.xlim(0, 1)
for i, v in enumerate(accuracy_rand):
    plt.text(v + 0.01, i, f"{v:.4f}", va='center', ha='left',
fontsize=10) # Adjust offset and precision as needed

plt.grid(axis='x', linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```



GridSearchCv

```
from sklearn.metrics import accuracy_score

# Calculate accuracy for each model
dt_accuracy_grid = accuracy_score(Yr_test, Y_pred_dt_grid)
lr_accuracy_grid = accuracy_score(Yr_test, Y_pred_lr_grid)

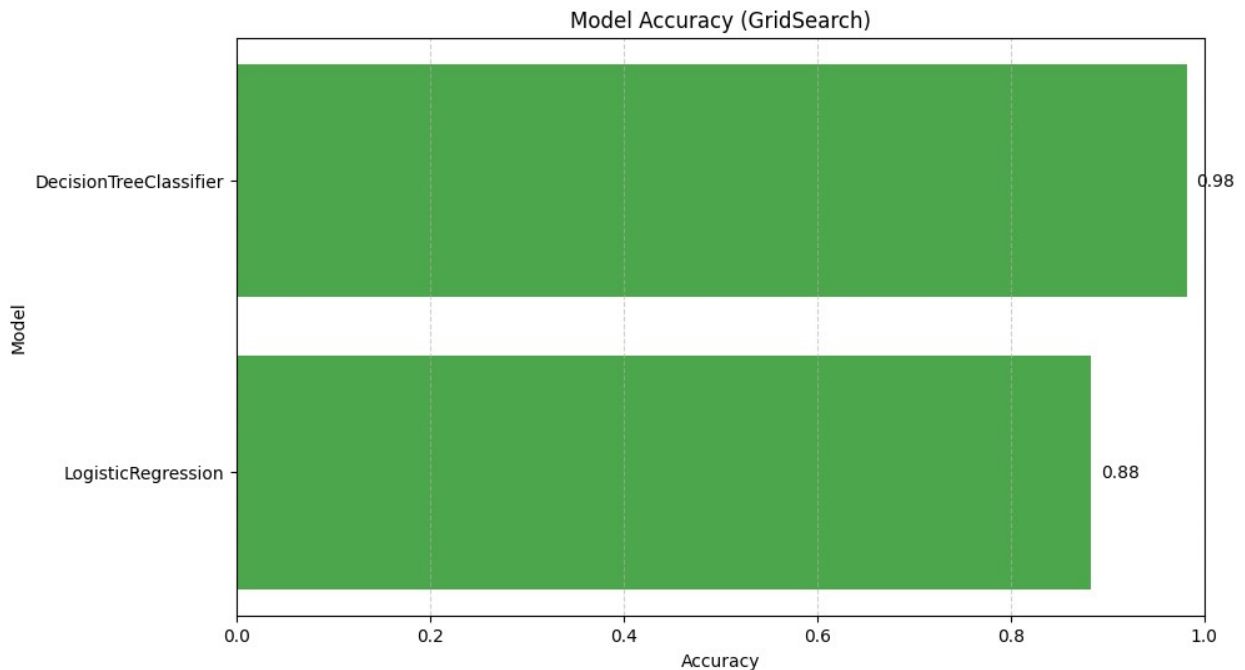
# Create the accuracy_rand list
accuracy_grid = [dt_accuracy_grid, lr_accuracy_grid]

# Model names list
model_names_grid = ['DecisionTreeClassifier', 'LogisticRegression']

import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))
plt.barh(model_names_grid, accuracy_grid, color='green', alpha=0.7)
plt.xlabel('Accuracy')
plt.ylabel('Model')
plt.title('Model Accuracy (GridSearch)')
plt.gca().invert_yaxis()
plt.xlim(0, 1)
for i, v in enumerate(accuracy_grid):
    plt.text(v + 0.01, i, f"{v:.2f}", va='center', ha='left',
    fontsize=10) # Adjust offset and precision as needed
```

```
plt.grid(axis='x', linestyle='--', alpha=0.6)
plt.show()
```



Without Hyperparameter tuning

```
from sklearn.metrics import accuracy_score

# Calculate accuracy for each model
dt_accuracy = accuracy_score(Yr_test, Y_pred_dt)
knn_accuracy = accuracy_score(Yr_test, Y_pred_knn)
lr_accuracy = accuracy_score(Yr_test, Y_pred_lr)
rf_accuracy = accuracy_score(Yr_test, Y_pred_rf)

# Create the accuracy_list
accuracy = [dt_accuracy, knn_accuracy, lr_accuracy, rf_accuracy]

# Model names list
model_names = ['DecisionTreeClassifier', 'KNN',
               'LogisticRegression', 'RandomForestClassifier']

import matplotlib.pyplot as plt

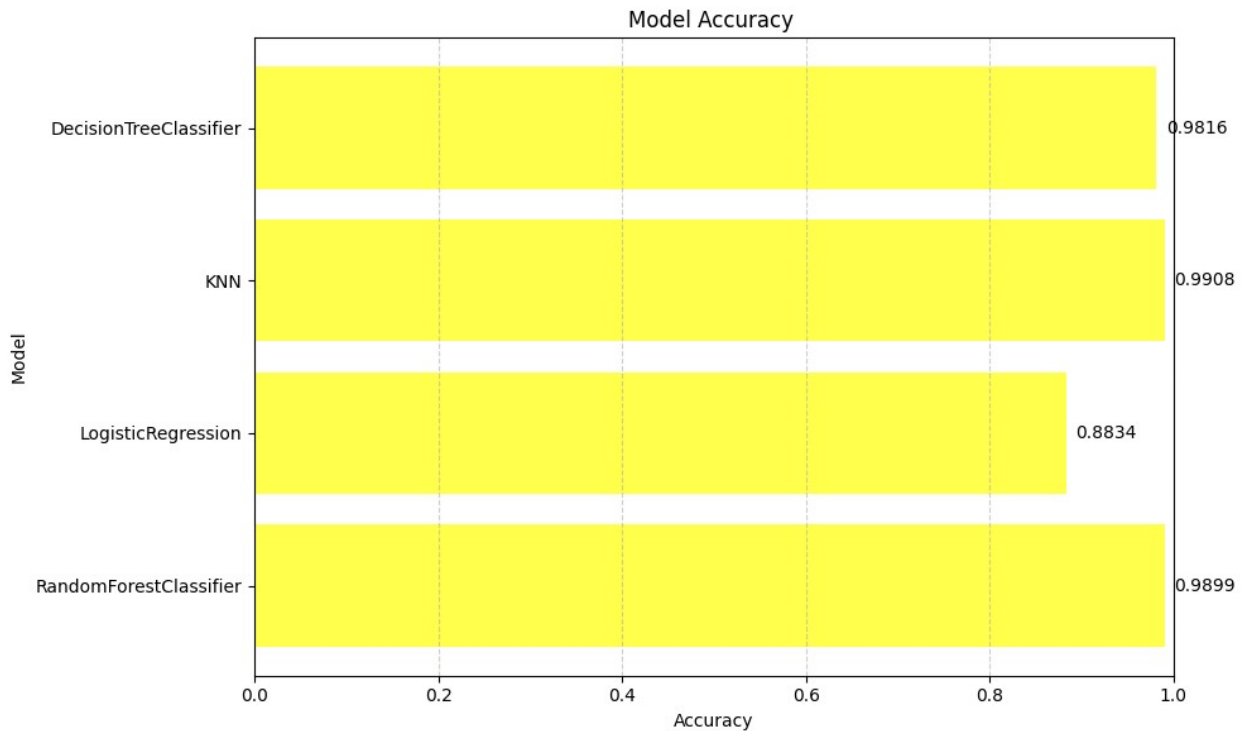
plt.figure(figsize=(10, 6))
plt.barh(model_names, accuracy, color='yellow', alpha=0.7)
plt.xlabel('Accuracy')
plt.ylabel('Model')
plt.title('Model Accuracy ')
plt.gca().invert_yaxis() # Invert y-axis for readability (highest
```

```

accuracy on top)
plt.xlim(0, 1)
for i, v in enumerate(accuracy):
    plt.text(v + 0.01, i, f"{v:.4f}", va='center', ha='left',
            fontsize=10) # Adjust offset and precision as needed

plt.grid(axis='x', linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()

```



BEST MODEL

```

# Combine the three dataframes
combined_metrics_df = pd.concat([metrics_df_rand, metrics_df_grid,
                                metrics_df])

# Find the model with the highest accuracy
best_model = combined_metrics_df['Accuracy'].idxmax()

print(f"Best model based on accuracy: {best_model}")

Best model based on accuracy: KNN_rand

```