

```

# NAME :- Akula Hema Venkata Sriram
# ROLL NO. :- 04
# REGISTRATION NO. :- 12210461
# SECTION :- K22BW
# COURSE CODE :- INT-254

#-----
# PROJECT TITLE :- CUSTOMER CHURN PREDICTION
#-----

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

telco_data = pd.read_csv("telco_customer_churn.csv")

# Display basic information about the dataset
print("Shape of the dataset:", telco_data.shape)
print("\nColumns in the dataset:")
print(telco_data.columns)
print("\nSample data:")
print(telco_data.head())

# Summary statistics
print("\nSummary statistics:")
print(telco_data.describe())

# Check for missing values
print("\nMissing values:")
print(telco_data.isnull().sum())

print("\nColumn names : ")
print(telco_data.columns.values)

print("\nColumns Data Types : ")
print(telco_data.dtypes)
# Check for duplicate rows
print("\nDuplicate rows:", telco_data.duplicated().sum())

# Visualize the distribution of the target variable 'Churn'
plt.figure(figsize=(8, 6))
sns.countplot(x='Churn', data=telco_data)
plt.title('Distribution of Churn')
plt.show()

# Visualize the distribution of numerical features
numerical_features =
telco_data.select_dtypes(include=[np.number]).columns.tolist()
telco_data[numerical_features].hist(figsize=(12, 10))

```

```
plt.suptitle('Distribution of Numerical Features')
plt.show()

# Visualize the correlation matrix
plt.figure(figsize=(10, 8))
correlation_matrix = telco_data[numerical_features].corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
fmt=".2f")
plt.title('Correlation Matrix')
plt.show()
```

C:\Users\Hitech\AppData\Local\Temp\ipykernel_4944\4206063596.py:2:
DeprecationWarning:
Pyarrow will become a required dependency of pandas in the next major
release of pandas (pandas 3.0),
(to allow more performant data types, such as the Arrow string type,
and better interoperability with other libraries)
but was not found to be installed on your system.
If this would cause problems for you,
please provide us feedback at
<https://github.com/pandas-dev/pandas/issues/54466>

```
import pandas as pd
```

Shape of the dataset: (7043, 21)

Columns in the dataset:

```
Index(['customerID', 'gender', 'SeniorCitizen', 'Partner',
'Dependents',
      'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
      'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',
'TechSupport',
      'StreamingTV', 'StreamingMovies', 'Contract',
'PaperlessBilling',
      'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'],
      dtype='object')
```

Sample data:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure
PhoneService \						
0	7590-VHVEG	Female	0	Yes	No	1
No						
1	5575-GNVDE	Male	0	No	No	34
Yes						
2	3668-QPYBK	Male	0	No	No	2
Yes						

3	7795-CFOCW	Male	0	No	No	45
No						
4	9237-HQITU	Female	0	No	No	2
Yes						

MultipleLines		InternetService	OnlineSecurity	...
DeviceProtection \				
0	No phone service	DSL	No	...
No				
1	No	DSL	Yes	...
Yes				
2	No	DSL	Yes	...
No				
3	No phone service	DSL	Yes	...
Yes				
4	No	Fiber optic	No	...
No				

TechSupport	StreamingTV	StreamingMovies	Contract
PaperlessBilling \			
0	No	No	No Month-to-month
Yes			
1	No	No	No One year
No			
2	No	No	No Month-to-month
Yes			
3	Yes	No	No One year
No			
4	No	No	No Month-to-month
Yes			

	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	Electronic check	29.85	29.85	No
1	Mailed check	56.95	1889.5	No
2	Mailed check	53.85	108.15	Yes
3	Bank transfer (automatic)	42.30	1840.75	No
4	Electronic check	70.70	151.65	Yes

[5 rows x 21 columns]

Summary statistics:

	SeniorCitizen	tenure	MonthlyCharges
count	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692
std	0.368612	24.559481	30.090047
min	0.000000	0.000000	18.250000
25%	0.000000	9.000000	35.500000
50%	0.000000	29.000000	70.350000
75%	0.000000	55.000000	89.850000
max	1.000000	72.000000	118.750000

Missing values:

customerID	0
gender	0
SeniorCitizen	0
Partner	0
Dependents	0
tenure	0
PhoneService	0
MultipleLines	0
InternetService	0
OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	0
StreamingTV	0
StreamingMovies	0
Contract	0
PaperlessBilling	0
PaymentMethod	0
MonthlyCharges	0
TotalCharges	0
Churn	0

dtype: int64

Column names :

```
['customerID' 'gender' 'SeniorCitizen' 'Partner' 'Dependents' 'tenure'
 'PhoneService' 'MultipleLines' 'InternetService' 'OnlineSecurity'
 'OnlineBackup' 'DeviceProtection' 'TechSupport' 'StreamingTV'
 'StreamingMovies' 'Contract' 'PaperlessBilling' 'PaymentMethod'
 'MonthlyCharges' 'TotalCharges' 'Churn']
```


Columns Data Types :

customerID	object
gender	object
SeniorCitizen	int64
Partner	object
Dependents	object
tenure	int64
PhoneService	object
MultipleLines	object
InternetService	object
OnlineSecurity	object
OnlineBackup	object
DeviceProtection	object
TechSupport	object
StreamingTV	object
StreamingMovies	object
Contract	object
PaperlessBilling	object




```
PaymentMethod      object
MonthlyCharges     float64
TotalCharges       object
Churn              object
```


```
dtype: object
```

```
Duplicate rows: 0
```



Distribution of Numerical Features





```
print(100*telco_data['Churn'].value_counts()/len(telco_data['Churn']))
print(telco_data['Churn'].value_counts())
telco_base_data=telco_data.copy()

Churn
No      73.463013
Yes     26.536987
Name: count, dtype: float64
Churn
No      5174
Yes     1869
Name: count, dtype: int64

telco_data.TotalCharges=pd.to_numeric(telco_data.TotalCharges,
errors='coerce')
telco_data.isnull().sum()
```

```

customerID      0
gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    11
Churn           0
dtype: int64

```

```

telco_data.loc[telco_data['TotalCharges'].isnull()==True]
telco_data.dropna(how='any',inplace=True)
telco_data.shape

```

```

(7032, 21)

```

```

print(telco_data['tenure'].max())

```

```

72

```

```

labels=["{0} - {1}".format(i,i+11)for i in range(1,72,12)]
telco_data['tenure_group']=pd.cut(telco_data.tenure,range(1,80,12),
right=False, labels=labels)

```

```

telco_data['tenure_group'].value_counts()

```

```

tenure_group
1 - 12      2175
61 - 72     1407
13 - 24     1024
25 - 36      832
49 - 60      832
37 - 48      762

```

```

Name: count, dtype: int64

```

```

telco_data.drop(columns=['customerID'],axis=1,inplace=True)
telco_data.head()

```


	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	\
0	Female	0	Yes	No	1	No	
1	Male	0	No	No	34	Yes	
2	Male	0	No	No	2	Yes	
3	Male	0	No	No	45	No	
4	Female	0	No	No	2	Yes	

	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	...	\
0	No phone service	DSL	No	Yes	...	
1	No	DSL	Yes	No	...	
2	No	DSL	Yes	Yes	...	
3	No phone service	DSL	Yes	No	...	
4	No	Fiber optic	No	No	...	

	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling	\
0	No	No	No	Month-to-month	Yes	
1	No	No	No	One year	No	
2	No	No	No	Month-to-month	Yes	
3	Yes	No	No	One year	No	
4	No	No	No	Month-to-month	Yes	

	PaymentMethod	MonthlyCharges	TotalCharges	Churn	tenure_group
0	Electronic check	29.85	29.85	No	1 - 12
1	Mailed check	56.95	1889.50	No	25 - 36
2	Mailed check	53.85	108.15	Yes	1 - 12
3	Bank transfer (automatic)	42.30	1840.75	No	37 - 48
4	Electronic check	70.70	151.65	Yes	1 - 12

[5 rows x 21 columns]

telco_data.head()

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	\
0	Female	0	Yes	No	1	No	
1	Male	0	No	No	34	Yes	
2	Male	0	No	No	2	Yes	
3	Male	0	No	No	45	No	
4	Female	0	No	No	2	Yes	

	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	...	\
0	No phone service	DSL	No	Yes	...	
1	No	DSL	Yes	No	...	
2	No	DSL	Yes	Yes	...	
3	No phone service	DSL	Yes	No	...	
4	No	Fiber optic	No	No	...	


	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling	\
0	No	No	No	Month-to-month	Yes	
1	No	No	No	One year	No	
2	No	No	No	Month-to-month	Yes	
3	Yes	No	No	One year	No	
4	No	No	No	Month-to-month	Yes	


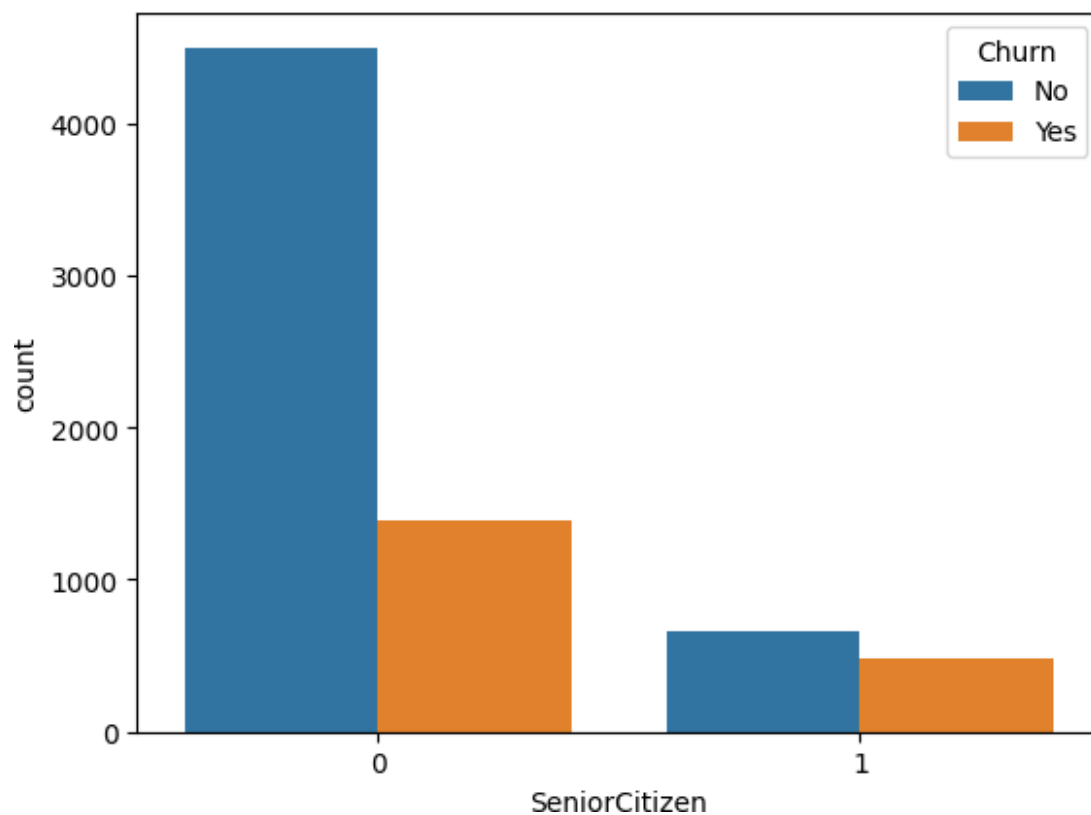
	PaymentMethod	MonthlyCharges	TotalCharges	Churn	tenure_group
0	Electronic check	29.85	29.85	No	1 - 12
1	Mailed check	56.95	1889.50	No	25 - 36
2	Mailed check	53.85	108.15	Yes	1 - 12
3	Bank transfer (automatic)	42.30	1840.75	No	37 - 48
4	Electronic check	70.70	151.65	Yes	1 - 12



[5 rows x 21 columns]

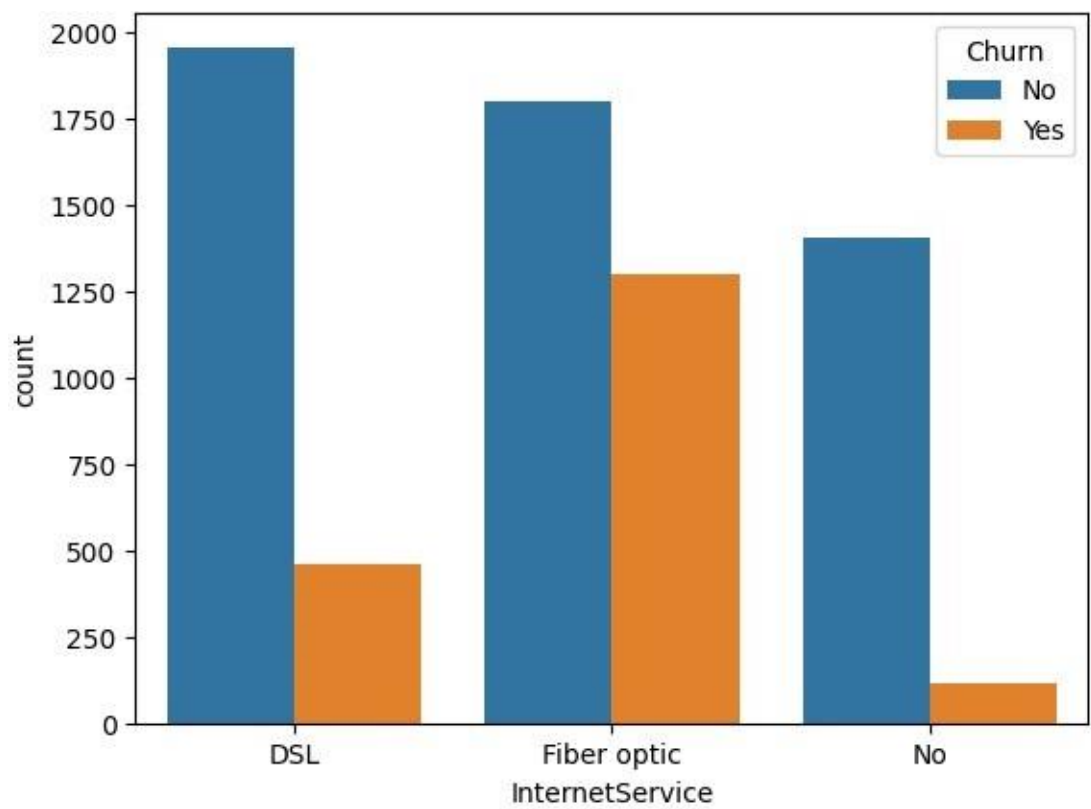

```
for i, predictor in
enumerate(telco_data.drop(columns=['Churn', 'TotalCharges', 'MonthlyChar
```



```
ges', 'tenure']])):  
plt.figure(i)  
sns.countplot(data=telco_data, x=predictor, hue='Churn')
```





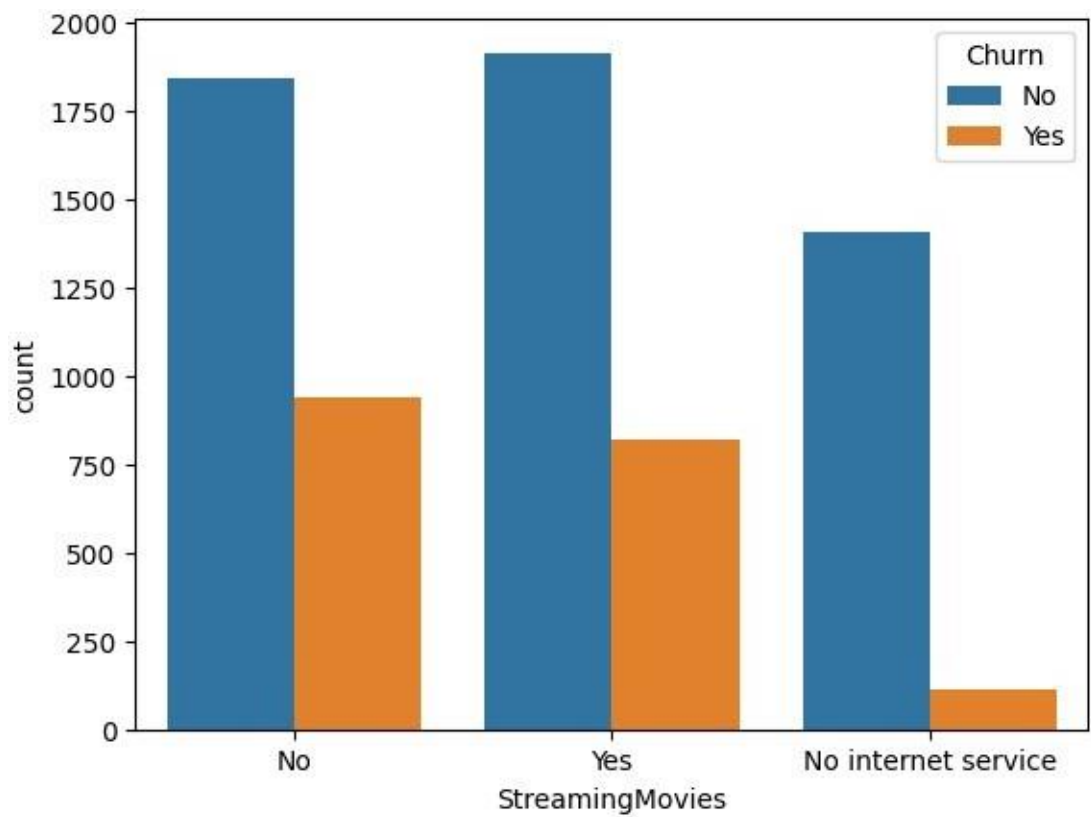



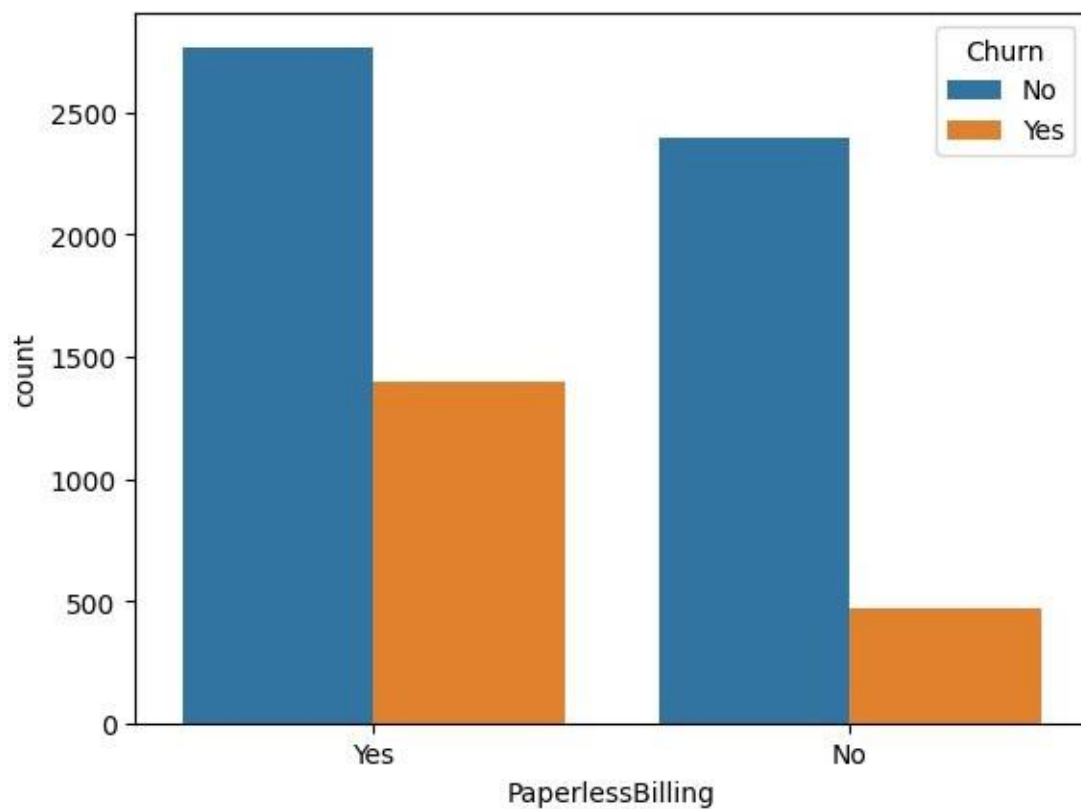

















```
telco_data['Churn']=np.where(telco_data.Churn=='Yes',1,0)
```

```
telco_data.head()
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	\
0	Female	0	Yes	No	1	No	
1	Male	0	No	No	34	Yes	
2	Male	0	No	No	2	Yes	
3	Male	0	No	No	45	No	
4	Female	0	No	No	2	Yes	

	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	...	\
0	No phone service	DSL	No	Yes	...	
1	No	DSL	Yes	No	...	
2	No	DSL	Yes	Yes	...	
3	No phone service	DSL	Yes	No	...	
4	No	Fiber optic	No	No	...	

	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling	\
0	No	No	No	Month-to-month	Yes	
1	No	No	No	One year	No	
2	No	No	No	Month-to-month	Yes	
3	Yes	No	No	One year	No	
4	No	No	No	Month-to-month	Yes	

	tenure_group	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0		Electronic check	29.85	29.85	0
1	1 - 12	Mailed check	56.95	1889.50	0
2	25 - 36	Mailed check	53.85	108.15	1
3	1 - 12	Bank transfer (automatic)	42.30	1840.75	0
4	37 - 48	Electronic check	70.70	151.65	1

```
[5 rows x 21 columns]
```

```
telco_data_dummies=pd.get_dummies(telco_data, dtype=int)
telco_data_dummies.head()
```

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges	Churn
gender_Female \					
0	0	1	29.85	29.85	0
1					
1	0	34	56.95	1889.50	0
0					
2	0	2	53.85	108.15	1
0					
3	0	45	42.30	1840.75	0
0					
4	0	2	70.70	151.65	1
1					

	gender_Male	Partner_No	Partner_Yes	Dependents_No	...	\
0	0	0	1	1	...	
1	1	1	0	1	...	
2	1	1	0	1	...	
3	1	1	0	1	...	
4	0	1	0	1	...	

	PaymentMethod_Bank transfer (automatic)	\
0	0	
1	0	
2	0	
3	1	
4	0	

	PaymentMethod_Credit card (automatic)	PaymentMethod_Electronic check	\
0	0		
1			
1	0		
0			
2	0		
0			
3	0		
0			
4	0		
1			

	PaymentMethod_Mailed check	tenure_group_1 - 12	tenure_group_13 - 24	\
0	0	1		
0				
1	1	0		
0				
2	1	1		

```

0
3
0
4
0

tenure_group_25 - 36 tenure_group_37 - 48 tenure_group_49 - 60 \
0 0 0
1 1 0
2 0 0
3 0 1
4 0 0


tenure_group_61 - 72
0 0
1 0
2 0
3 0
4 0

[5 rows x 52 columns]

plt.figure(figsize=(20,8))
telco_data_dummies.corr()
['Churn'].sort_values(ascending=False).plot(kind='bar')


<Axes: >

```



```
plt.figure(figsize=(12,12))
sns.heatmap(telco_data_dummies.corr(),cmap="Paired")
```

<Axes: >



```
df_target_churn0=telco_data.loc[telco_data["Churn"]==0]
df_target_churn1=telco_data.loc[telco_data["Churn"]==1]
```

#model

```
import pandas as pd
from sklearn import metrics
from sklearn.model_selection import train_test_split
```

```
df=telco_data_dummies.copy()
df.head()
```

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges	Churn
gender_Female \					
0	0	1	29.85	29.85	0
1					
1	0	34	56.95	1889.50	0
0					
2	0	2	53.85	108.15	1
0					
3	0	45	42.30	1840.75	0
0					
4	0	2	70.70	151.65	1
1					

	gender_Male	Partner_No	Partner_Yes	Dependents_No	...	\
0	0	0	1	1	...	
1	1	1	0	1	...	
2	1	1	0	1	...	
3	1	1	0	1	...	
4	0	1	0	1	...	

	PaymentMethod_Bank transfer (automatic)	\
0	0	
1	0	
2	0	
3	1	
4	0	

	PaymentMethod_Credit card (automatic)	PaymentMethod_Electronic
check \		
0	0	
1		
1	0	
0		
2	0	
0		
3	0	
0		
4	0	
1		

	PaymentMethod_Mailed check	tenure_group_1 - 12	tenure_group_13 -
24 \			
0	0	1	

0		
1	1	0
0		
2	1	1
0		
3	0	0
0		
4	0	1
0		

	tenure_group_25 - 36	tenure_group_37 - 48	tenure_group_49 - 60 \
0	0	0	0
1	1	0	0
2	0	0	0
3	0	1	0
4	0	0	0

	tenure_group_61 - 72
0	0
1	0
2	0
3	0
4	0

[5 rows x 52 columns]

```
X=df.drop("Churn",axis=1)
Y=df['Churn']
```

```
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.2,random_state=42)
```

```
from sklearn.tree import DecisionTreeClassifier
dt=DecisionTreeClassifier(criterion='gini',random_state=100,max_depth=6,min_samples_leaf=8)
dt.fit(X_train,Y_train)
```

```
DecisionTreeClassifier(max_depth=6, min_samples_leaf=8,
random_state=100)
```

```
Y_predict=dt.predict(X_test)
```

```
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
print(classification_report(Y_test,Y_predict))
print(confusion_matrix(Y_test,Y_predict))
```

	precision	recall	f1-score	support
0	0.84	0.85	0.85	1033
1	0.58	0.56	0.57	374

accuracy			0.78	1407
macro avg	0.71	0.71	0.71	1407
weighted avg	0.77	0.78	0.78	1407

```
[[883 150]
 [164 210]]
```

```
from sklearn.ensemble import RandomForestClassifier
rf=RandomForestClassifier(n_estimators=100)
rf.fit(X_train,Y_train)
Y_rf_pred=rf.predict(X_test)
print(classification_report(Y_test,Y_rf_pred))
print(confusion_matrix(Y_test,Y_rf_pred))
```

	precision	recall	f1-score	support
0	0.82	0.89	0.85	1033
1	0.59	0.45	0.51	374

accuracy			0.77	1407
macro avg	0.71	0.67	0.68	1407
weighted avg	0.76	0.77	0.76	1407

```
[[918 115]
 [206 168]]
```

```
from imblearn.combine import SMOTEENN
sm=SMOTEENN()
X_resampled,Y_resampled=sm.fit_resample(X,Y)
Xr_train,Xr_test,Yr_train,Yr_test=train_test_split(X_resampled,Y_resampled,test_size=0.2,random_state=42)
print(Y_resampled.value_counts())
print(Y.value_counts())
```

```
Churn
1    3196
0    2652
Name: count, dtype: int64
```

```
Churn
0    5163
1    1869
Name: count, dtype: int64
```

```
from sklearn.tree import DecisionTreeClassifier
dt=DecisionTreeClassifier(criterion='gini',random_state=100,max_depth=6,min_samples_leaf=8)
dt.fit(Xr_train,Yr_train)
Yr_predict=dt.predict(Xr_test)
print(classification_report(Yr_test,Yr_predict))
print(confusion_matrix(Yr_test,Yr_predict))
```

	precision	recall	f1-score	support
0	0.92	0.94	0.93	556
1	0.94	0.93	0.94	614
accuracy			0.93	1170
macro avg	0.93	0.93	0.93	1170
weighted avg	0.93	0.93	0.93	1170

```
[[520  36]
 [ 43 571]]
```

```
from sklearn.ensemble import RandomForestClassifier
rf=RandomForestClassifier(n_estimators=100)
rf.fit(Xr_train,Yr_train)
Yr_rf_pred=rf.predict(Xr_test)
print(classification_report(Yr_test,Yr_rf_pred))
print(confusion_matrix(Yr_test,Yr_rf_pred))
```

	precision	recall	f1-score	support
0	0.97	0.96	0.96	556
1	0.96	0.97	0.97	614
accuracy			0.96	1170
macro avg	0.96	0.96	0.96	1170
weighted avg	0.96	0.96	0.96	1170

```
[[532  24]
 [ 18 596]]
```

```
from sklearn.neural_network import MLPClassifier

# Model 6: Neural Network Classifier
print("\nModel 6: Neural Network Classifier")
model_nn = MLPClassifier(hidden_layer_sizes=(100,), max_iter=1000)
model_nn.fit(Xr_train, Yr_train)
yr_pred_nn = model_nn.predict(Xr_test)
print("Classification Report:")
print(classification_report(Yr_test, yr_pred_nn))
print("Confusion Matrix:")
print(confusion_matrix(Yr_test, yr_pred_nn))
```

Model 6: Neural Network Classifier
Classification Report:

	precision	recall	f1-score	support
0	0.93	0.94	0.94	556
1	0.95	0.94	0.94	614

accuracy			0.94	1170
macro avg	0.94	0.94	0.94	1170
weighted avg	0.94	0.94	0.94	1170

Confusion Matrix:

```
[[523  33]
 [ 37 577]]
```

```
from sklearn.model_selection import GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, confusion_matrix
```

Define pipeline

```
pipe_dt = Pipeline([
    ('clf', DecisionTreeClassifier())
])
```

Define parameter grid

```
param_grid_dt = {
    'clf__criterion': ['gini', 'entropy'],
    'clf__max_depth': [None, 10, 20, 30, 40, 50],
    'clf__min_samples_split': [2, 5, 10],
    'clf__min_samples_leaf': [1, 2, 4],
    'clf__max_features': ['sqrt', 'log2', None]
}
```

Perform GridSearchCV

```
grid_dt = GridSearchCV(pipe_dt, param_grid_dt, cv=5)
grid_dt.fit(Xr_train, Yr_train)
```

Print best parameters

```
print("Best Parameters (GridSearchCV):", grid_dt.best_params_)
```

Predict on the testing set using the best model

```
best_classifier_dt = grid_dt.best_estimator_
yrr_pred_dt = best_classifier_dt.predict(Xr_test)
```

Evaluate the model

```
print("\nClassification Report:")
print(classification_report(Yr_test, yrr_pred_dt))
```

```
print("\nConfusion Matrix:")
```

```
print(confusion_matrix(Yr_test, yrr_pred_dt))
```

```
Best Parameters (GridSearchCV): {'clf__criterion': 'gini',
'clf__max_depth': 10, 'clf__max_features': None,
'clf__min_samples_leaf': 2, 'clf__min_samples_split': 10}
```

Classification Report:

precision	recall	f1-score	support
-----------	--------	----------	---------

0	0.93	0.94	0.93	556
1	0.94	0.93	0.94	614
accuracy			0.94	1170
macro avg	0.94	0.94	0.94	1170
weighted avg	0.94	0.94	0.94	1170

Confusion Matrix:

```
[[522  34]
 [ 41 573]]
```

```
from imblearn.over_sampling import SMOTE
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report
from sklearn.model_selection import GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder

# Initialize and train Random Forest classifier
rf_clf = RandomForestClassifier(random_state=42)
param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 5, 10],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}
grid_search = GridSearchCV(rf_clf, param_grid, cv=5,
scoring='accuracy')
grid_search.fit(Xr_train, Yr_train)

# Get the best estimator
best_rf_clf = grid_search.best_estimator_

# Evaluate on test data
Yrgd_pred = best_rf_clf.predict(Xr_test)
print(classification_report(Yr_test, Yrgd_pred))
```

	precision	recall	f1-score	support
0	0.97	0.96	0.96	556
1	0.96	0.97	0.97	614
accuracy			0.96	1170
macro avg	0.96	0.96	0.96	1170
weighted avg	0.96	0.96	0.96	1170

```

from xgboost import XGBClassifier

# Model 7: XGBoost Classifier
print("\nModel 7: XGBoost Classifier")
model_xgb = XGBClassifier()
model_xgb.fit(Xr_train, Yr_train)
yr_pred_xgb = model_xgb.predict(Xr_test)
print("Classification Report:")
print(classification_report(Yr_test, yr_pred_xgb))
print("Confusion Matrix:")
print(confusion_matrix(Yr_test, yr_pred_xgb))

```

Model 7: XGBoost Classifier

Classification Report:

	precision	recall	f1-score	support
0	0.96	0.97	0.96	556
1	0.97	0.96	0.97	614
accuracy			0.97	1170
macro avg	0.97	0.97	0.97	1170
weighted avg	0.97	0.97	0.97	1170

Confusion Matrix:

```

[[538  18]
 [ 22 592]]

```